

CHAPTER 5

AI Safety and Operational Safeguards for Diplomacy

From principle to practice: making AI safe in the daily work of statecraft

The question for foreign ministries is no longer whether artificial intelligence will be used in diplomatic work, but whether its use will be careful, checked and accountable. Safety is not something added on after a system goes live; it is the discipline that lets AI widen a diplomat's judgement without ever replacing it.

Introduction

Earlier chapters set out what artificial intelligence can do for diplomacy, what can go wrong, and where the technology falls short. Chapter 1 described intelligent statecraft as AI helping diplomats see more, not replacing their judgement — human responsibility for interpreting and representing the state stays intact. Chapters 2 and 3 set out, in turn, the risks AI brings into diplomatic work and the limits it cannot reliably overcome. Chapter 4 placed these concerns inside a framework of ethics and accountability.

This chapter moves from describing problems to building solutions. It answers the practical question every ministry must face before adopting AI faster than it can govern it: **what concrete safeguards make AI safe to use in diplomatic work?** The focus here is deliberately hands-on. Earlier chapters described where harm can arise and what principles should guide use; this chapter describes the steps that connect the two.

Before setting out those steps, it helps to draw a distinction between a few key terms that institutional policy discussions often blur.


- Cybersecurity protects systems and networks from intrusion.
- Information security protects the confidentiality, integrity and availability of data.
- AI safety is about how an AI system behaves internally and how to ensure it behaves reliably and doesn't cause harm, even when they're powerful, widely deployed, or operating in unexpected situations.

A model can be completely secure against outside attack and still produce a confident, well-written and wrong diplomatic brief. The International AI Safety Report 2026 warns that advanced AI can make existing harms worse, and that some of its most serious failures only become visible once it is used in the real world (International AI Safety Report, 2026). Shah et al. (2025) draw a similar line between safety — preventing unintended harm that comes from the system itself — and security — preventing deliberate misuse by an outside attacker. **Foreign ministries and diplomacy need both, but each calls for different routines, different language and different people responsible for it.**

The thread running through all six sections of this chapter is Malaysia's emerging role in the Association of Southeast Asian Nations (ASEAN) AI-safety agenda. A foreign ministry that can show disciplined,

documented safeguards is better placed to help shape regional AI-safety norms than one that simply buys tools. The sections that follow show what that discipline looks like in practice.

Figure 5.1 The chapter's framework — AI safety for diplomatic statecraft



Source: Our synthesis from Shah et al. (2024), International AI Safety Report (2026).

5.1 AI Risk Assessment and Classification for Diplomatic Use Cases

Effective AI safety starts not with technology but with sorting risk by kind and size.

The National Institute of Standards and Technology (NIST) AI Risk Management Framework (AI RMF 1.0) treats risk management as an ongoing job — Govern, Map, Measure, Manage — whose first task is to understand risk in the specific setting where a system is used, rather than lean on general reassurances from a vendor (National Institute of Standards and Technology, 2023). For a foreign ministry, that first step decides, before any tool is adopted, how much care each use case needs.

One useful way to make this concrete is to look at every use case through the four layers introduced in Figure I.1. Risk does not sit in one place; it builds up across layers, and any classification that skips a layer is incomplete.

Table 5.1a The four-layer AI-enabled foreign-affairs system — classification lens

Layer	What can go wrong	Diplomatic example
Model	Hallucination, bias, or the model being confidently wrong in its own output.	A translation model confidently mistranslates a binding legal term.
Application	How the model is wired into work — finding documents, drafting, summarising.	A step that pulls in outside documents brings in a tampered embassy cable the model alone would never surface.
Institutional	Who approves, reviews and is accountable.	A consular brief accepted without a named reviewer.
Infrastructure	Where data and compute sit, under whose jurisdiction.	A vendor with silent update rights changes model behaviour after approval.

Source: Our synthesis.

Classification labels each use case by the highest layer where a serious failure could happen, then sets the level of review accordingly.

The nine-category list of diplomatic risks below gives the words for what kind of harm each layer might cause; the tier table (Table 5.1c) sets out how much care each level needs. **The ninth category — catastrophic — is why classification can never be a formality.** The International AI Safety Report 2026 treats catastrophic AI risk, including runaway escalation and loss of control, as a live policy concern rather than a distant scenario (International AI Safety Report, 2026).

Table 5.1b The nine kinds of diplomatic AI risk

#	Category	Diplomatic example of what can go wrong
1	Epistemic	Confident hallucination — a fluent briefing note that quietly states a wrong figure.
2	Linguistic-cultural	Lost nuance or culturally tone-deaf phrasing that offends or misleads a counterpart.
3	Information-integrity	Fabricated sources, deepfakes or manipulated evidence entering the diplomatic record.
4	Cybersecurity	A hidden instruction smuggled into a document, or tampered data pulled from a connected classified system.
5	Decision	AI distorting or crowding out human judgement in a high-stakes call.
6	Institutional	Unclear ownership or missing review when something breaks.
7	Sovereignty	Data or jurisdiction held offshore; vendor control of silent model updates.
8	Systemic	Errors spreading across missions and partners; gradual erosion of institutional trust.

#	Category	Diplomatic example of what can go wrong
9	Catastrophic	A single failure cascading to bilateral rupture, conflict escalation or loss of life.

Source: Authors' synthesis from International AI Safety Report (2026) and NIST AI RMF 1.0 (2023).

Table 5.1c Risk classification matrix for diplomatic AI use cases

Tier	Example use case	Sensitivity	Consequence if wrong	Required review
Low	Summarising open-source news; meeting-note tidy-up	Low	Minimal	Self-check; spot review
Moderate	Research synthesis for a briefing; first-draft public explainer	Medium	Reputational; minor policy	Named desk officer
High	Negotiation position support; consular advice; counterpart risk assessment	High	Bilateral; legal; personal safety	Senior review + source verification + approval chain
Critical	Autonomous action on classified material; external national-position commitment	Highest	National interest; sovereignty	Executive sign-off; legal; recorded authority

Source: Our synthesis from NIST AI RMF 1.0 (2023) and our survey and interview data.

Our survey of 64 diplomatic officers shows that officers already think in layers without using that word for it. Forty-six respondents raised concerns about confidentiality or data leaking out on their own — an instinctive worry about the infrastructure layer — while 15 pointed to high-stakes decisions and negotiations, a concern about the institutional layer (Survey for Diplomats, 2026). **We classify risk precisely so that this instinct can be checked and acted on, not just felt.**

The International AI Safety Report 2026 explains why judging a system by the model alone is not enough: advanced AI can make existing harms worse, and its failures often only show up in real use, not on a vendor's benchmark (International AI Safety Report, 2026). This is exactly why the four-layer view matters — **a model that passes a vendor's test can still fail at the application, institutional or infrastructure layer once it is actually deployed.**

Hence, the catastrophic category is not a hypothetical. Consider a scenario built on the synthetic-identity breach discussed in Chapter 2: a convincing AI-generated message, made to look like it came from a senior official, commits a government to a negotiating position it never agreed to. If this is caught too late, such an output could harden positions in an ongoing dispute, put a separate negotiation at risk, or — in the worst case — help trigger an escalation that is hard to reverse. Classifying

this kind of use case at the Critical tier, with sign-off required from the top, is the administrative safeguard built to stop exactly that chain of events.

What you can do today (5.1) — model and institutional layers
Before adopting a new AI tool, walk it through all nine risk categories in Table 5.1b in under five minutes and circle any that clearly do not apply. What is left uncircled is the ministry's genuine exposure — treat that as the classification, not the vendor's claimed safety.

5.2 Evaluation, Validation and Continuous Monitoring

Classification tells a ministry how much care a use case needs. Evaluation and validation check whether a system actually deserves that trust — and, as the four-layer model makes clear, the answer is different at every layer. Shevlane et al. (2023) argue that testing a model well before use is critical for catching extreme risks, because developers must find dangerous capabilities before deployment, and models can show new abilities that even their own creators did not expect.

For diplomacy, the lesson is simple: a tool should not enter sensitive work just because a vendor says it is safe. Evaluating these models must happen at every layer, and the NIST Generative AI Profile — the companion document to the AI RMF 1.0 — supports this by treating evaluation, incident response and ongoing monitoring as parts of one continuing risk-management effort, not one-off checks (National Institute of Standards and Technology, 2024). A model that performs well in a simple chat window may behave very differently once it is connected to classified files, email, document systems or outside communication channels; it is the whole system in use that needs testing, not the model on its own.

Much of this testing now runs on standardised and automated benchmarks rather than one-off and manual. A benchmark is a fixed set of tasks used to score any model the same way, so results can be compared across tools and over time. MLCommons' AILuminate benchmark, for example, grades a chatbot's resistance to twelve categories of harmful prompts on a simple five-tier scale from Poor to Excellent, and the UK AI Security Institute's open-source Inspect Evals repository gives any team more than two hundred ready-made evaluations covering coding, reasoning, cybersecurity and safety behaviour, so a ministry does not have to build a test suite from nothing.

These are useful starting points, not a substitute for testing on the ministry's own documents and languages — and a good score today is no guarantee for tomorrow. Akhtar et al. (2026) find that nearly half of sixty widely used benchmarks have "saturated," meaning leading models now bunch so close to the ceiling score that the benchmark can no longer tell them apart, so a headline benchmark result should always be checked against how recently it was set and how it holds up outside the lab.

Table 5.2 What to evaluate at each layer

Layer	What to evaluate
Model	Ability and accuracy on real diplomatic tasks, including hallucination and overconfidence, using both standard benchmarks and diplomacy-specific tests. Shevlane's "dangerous-capability evaluation" belongs here.
Application	How well it finds the right documents, translation accuracy, draft coherence, using the ministry's own documents. A model that scores well on a benchmark can still fail the moment it touches a cable database.
Institutional	Whether officers notice errors, challenge recommendations and keep the skills to work without the tool — the check on how people actually use the system that Ganguli et al. (2022) and the International AI Safety Report 2026 treat as essential.
Infrastructure	Continuity: does the workflow survive a vendor outage, a silent model update or a jurisdiction change? Monitoring here is often the only warning that a system has drifted.

Source: Our synthesis from Shevlane et al. (2023), Ganguli et al. (2022), NIST AI RMF 1.0 (2023).

Sidhu et al. (2026) extend this idea to what happens after deployment: AI systems can fail in ways that pre-launch testing, including benchmarks, never predicted, and managing AI incidents well requires consistent definitions, categories, monitoring habits and reporting channels. Their research finds that existing frameworks define and sort incidents inconsistently — a gap a ministry closes by watching all four layers on an ongoing basis, rather than trusting a one-off benchmark certificate for the model.

Case studies — 5.2 Evaluation, Validation and Continuous Monitoring

Shevlane et al. (2023) show that dangerous capabilities can appear without the developer meaning them to — a model can gain an ability no benchmark ever checked for, which is exactly why model-layer evaluation for diplomatic use must actively look for the unexpected rather than tick boxes on an existing test. Our survey backs this up: 47 of 64 respondents raised accuracy, hallucination or double-checking worries without being asked (Survey for Diplomats, 2026), showing that officers already do informal checking at the institutional layer and deserve a proper, formal version of it.

The Deloitte case raised in Chapter 3 is a good example at the application layer too. A consulting firm handed a government a report containing AI-written content that had not been checked against original sources; the report was withdrawn and the firm agreed to a refund (Tadros and Karp, 2025). No benchmark score would have caught this, because the failure was in how the tool's output was used, not in the model itself — testing the whole application, not just the model's benchmark result, must confirm that finding sources, writing summaries and crediting them correctly all work on the ministry's real documents before that tool is trusted with sensitive work.

Benchmarks put model-layer evaluation within reach even for ministries with limited technical staff. AILuminate's twelve-category scorecard and Inspect Evals' library of over two hundred open-source tests mean a small team can get a defensible first read on a new tool within a day, at almost

no cost (Ghosh et al., 2025; AI Security Institute, 2024), although, the benchmark result is a starting filter, not a final verdict — the closing test always has to be the tool's performance on the ministry's own material.

What you can do today (5.2) — model and application layers
Before trusting any AI output in a brief or note, apply one rule: find the original source for the key fact, or do not use it. Checking a single loaded fact against a trustworthy source turns passive reading into active checking at the model layer. Chapter 7 looks at how this habit is built into formal training and AI literacy across the service.

5.3 Red-Teaming and Adversarial Testing

Evaluation asks whether a system works as intended under normal conditions. Red-teaming asks whether it fails dangerously under pressure. Ganguli et al. (2022) describe red-teaming as using manual or automatic methods to deliberately push a language model towards harmful outputs, treating it as one important tool among several for reducing harm. In diplomacy, the danger is not only a technical glitch but deliberate misuse across the whole system.

One risk especially relevant to diplomatic use is an AI assistant overstepping its limits: a translator, summariser or drafter given access to tools, documents and workflows wide enough to let it act beyond what was ever approved.

Red-team scenarios for diplomatic systems include: a tampered briefing document that tricks the model into producing false statements attributed to someone; a hidden instruction buried in an outside report that the model reads and follows; a translation that quietly shifts a negotiating position; an AI agent that sends a message nobody authorised; and a vendor outage that removes analysis support during a regional crisis. Each of these is a different kind of failure that needs a different kind of test.

Table 5.3 *What to red-team at each layer*

Layer	What to probe
Model	Tricks that bypass its safety limits, biased or manipulable outputs, unexpected harmful abilities the vendor did not test for.
Application	Overstepping its limits — tampered briefing documents, hidden instructions in outside reports, translations that shift a negotiating position.
Institutional	The human and procedural line of defence — can a convincing fake sense of authority push an officer to skip review, or an agent to send an unauthorised message?
Infrastructure	Channels and continuity — can an attacker exploit a vendor's update path, or cut off analysis support during a crisis?

Source: Our synthesis from Ganguli et al. (2022).

Ganguli et al. (2022) describe a pattern seen in red-teaming: a model's resistance to being pushed towards harmful outputs changes unpredictably as its abilities grow.

This means a ministry's schedule for red-teaming has to keep pace with model updates rather than being treated as a one-off check — a direct concern at the infrastructure layer, since vendors often update models without notice. There are also open-source tools to do model-layer testing within reach even for ministries with limited technical staff. Tools such as [Garak](#) and Microsoft's [PyRIT](#) let a small team check a model for safety-bypass tricks and hidden-instruction attacks at almost no cost; running a downloadable model on ministry hardware through [Ollama](#) means diplomatic data never has to leave the building, addressing the infrastructure-layer sovereignty worry at the same time.

Case studies — 5.3 Red-Teaming and Adversarial Testing

The Rubio impersonation incident from Chapter 2 shows a red-teaming failure at the institutional layer: an attacker used AI to create convincing fake messages that appeared to come from a senior official, and these reached foreign counterparts before anyone could check them (Lee, 2025). A red-team exercise that specifically tested whether a fake sense of authority could push officers to skip the call-back check would have caught this weakness early and cheaply.

A live, Malaysia-specific case shows the same kind of failure at the application layer, at a far larger scale than any single traveller's mishap. In September 2026, Anthropic disclosed that it had disrupted a commercial election-manipulation platform, built on its own Claude models, that an Istanbul-based firm, BBS Bilisim Teknolojileri, had marketed to clients as a "military-grade, AI-driven, real-time political operations ecosystem" (Anthropic, 2026). The operators fed real Malaysian census and electoral data into Claude to profile voters constituency by constituency across all 222 parliamentary seats, targeting racial, religious and royal fault lines, and ran a network of roughly 1,000 fake social-media accounts alongside a fabricated news outlet, "Malaysia Pulse" (Anthropic, 2026; The Edge Malaysia, 2026). Tellingly, when Claude itself refused one request after flagging a document as material for political defamation, the operators did not give up; they simply rewrote the request in softer, sanitised wording and kept building the same capability (Anthropic, 2026; Koh, 2026).

That is a model-layer safeguard being negotiated around rather than defeated outright — exactly the kind of adversarial workaround a red-team exercise should probe for before deployment, not discover only once a vendor's own threat-intelligence team catches it. The Malaysian Communications and Multimedia Commission (MCMC) responded within a day of the report's publication, stating publicly that "activities of the nature described in the report do not form part of MCMC's functions or mandate" and that it was seeking further information before deciding its next steps (MCMC, 2026).

That response itself is a lesson for the institutional layer: **even a capable, engaged regulator can find that a specific harm falls outside its own mandate** — exactly the kind of accountability gap Chapter 6 argues foreign ministries must help close through clearer inter-agency lines, rather than assume

someone else already owns it. A red-team test that specifically tries to reformulate a refused request in softer language, then checks whether the safeguard still holds, would test exactly this kind of failure.

The International AI Safety Report 2026 treats this kind of testing as part of a wider culture of checking systems across countries — a habit Malaysia can bring into ASEAN security exercises rather than simply buy as a vendor service, building regional skill rather than dependence (International AI Safety Report, 2026).

What you can do today (5.3) — institutional and infrastructure layers
Before trusting any AI drafting, summarising or translation tool with real material, spend fifteen minutes trying to break it yourself: feed it a fake urgent instruction, a hidden note buried inside an attached document, or a softened rewrite of a request it already refused.
Free toolkits such as Garak and PyRIT — paired with a locally run model through Ollama — let one officer do this in an afternoon without sending sensitive material to a vendor. If the tool complies with any of them, treat that as a finding to fix, not a one-off glitch: Ganguli et al. (2022) note that a model's resistance to manipulation shifts unpredictably as its abilities grow, so the test needs repeating with every update.

5.4 Internal AI Guidelines, Escalation Protocols and Safe-Use Policies for Foreign Ministries

Classification, evaluation and red-teaming matter, but they are not enough unless they are written into clear internal rules and a working escalation path. Looked at through the four layers, the guideline must say something specific about each one.

Table 5.4 *What internal AI guidelines must specify at each layer*

Layer	What the guideline must specify
Model	Which models are approved for which tiers; a ban on unvetted models for sensitive work; any deployed model must carry a current evaluation record.
Application	What may be entered into a system, what must remain on approved platforms, and limits on pulling in outside documents and using extra tools. The Open Worldwide Application Security Project (OWASP) Gen AI Top 10 categories — hidden-instruction attacks, risks from outside software or vendors, giving the system too much freedom to act, data leakage — map here (OWASP, 2024).
Institutional	Named human review per tier; who approves tools at missions; and the escalation protocol — who is notified, through what secure channel, on what timeline, with what action taken to contain it. The NIST AI RMF Govern function makes accountable, transparent use an organisational requirement (National Institute of Standards and Technology, 2023).

Layer	What the guideline must specify
Infrastructure	Rules on where data is stored and whose laws apply; requirements that vendors give notice before updates; plans for staying operational and for leaving a vendor if needed — the sovereignty dimension the guideline must not leave out.

Source: Authors' synthesis from NIST AI RMF 1.0 (2023), OWASP Gen AI Top 10 (2024), and the Ministry of Science, Technology and Innovation (MOSTI) National Guidelines on AI Governance and Ethics (2024).

Our survey points to a readiness gap these guidelines must close: only 15 of 64 respondents agreed that the rules for using AI were clear and consistent while 31 disagreed (Survey for Diplomats, 2026).

The MOSTI National Guidelines on AI Governance and Ethics set out national principles — privacy, security, transparency, accountability — but these still need to be translated into working rules for foreign affairs & diplomacy (Ministry of Science, Technology and Innovation Malaysia, 2024). This also requires a similar action from us for the Public Sector AI Adoption Guidelines as we need data classification and vendor disclosure at agency level to be implemented; how both apply to foreign affairs and diplomacy is covered in Chapter 6.

The guideline and escalation path work best when built around a diplomatic AI safety case: a clear, evidence-backed argument that a specific use of AI is safe enough to run.

For each significant use case, the ministry can start to write down the intended purpose, banned uses, who uses it and who is affected, how sensitive the data is, which model and vendor, what tools and permissions it has, its known limits, its test and red-team results, where a human must decide, what happens if something goes wrong, the fallback plan, and the conditions under which it must be switched off. **This turns "responsible AI" from a slogan into something that can be checked and approved.**

Our interviews also show that guidance is only as strong as the training behind it. Ambassador Zamshari Shaharan put it plainly: awareness cannot be built just by announcing that AI exists; officers need to understand hallucination, prompting, the limits of sources, and where personal use ends and official use begins, and for that "you need to do training." Training and guidelines are two sides of the same need — the institution has to make the safe choice the easiest one.

Figure 5.4 Escalation-protocol flow for AI incidents in diplomatic work




Figure 5.4 Escalation-protocol flow for AI incidents in diplomatic work

Source: Our synthesis from Sidhu & Scholefield (2026), NIST AI RMF 1.0 (2023), OWASP Gen AI Top 10 (2024).

Case studies — 5.4 Internal AI Guidelines, Escalation Protocols and Safe-Use Policies

Sidhu and Scholefield (2026) find that definitions of "incident" differ across regulators worldwide, which means our safety-case template is not paperwork for its own sake but the only reliable record of what was tested, by whom, and with what result. It also can become the seed of the regional incident database that Chapters 6 and 7 will talk more on ASEAN-wide cooperation.

Given that many enterprises will delay a quarter of their planned 2026 AI spending into 2027 because the financial payoff has been difficult to demonstrate (Gross, 2025) and other 2026 surveys have found only a small share of enterprises reporting measurable, at-scale returns despite heavy annual AI budgets, hence, this shows that: **a safety case does more than manage risk: it is also the paperwork that lets a ministry show, in plain terms, what a given AI use actually achieved — which is exactly the evidence procurement and finance officers need before committing further budget to a tool.**

What you can do today (5.4) — institutional layer

Write down, in one sentence, who you would call and by what channel if an AI-assisted output looked wrong — then check that answer against Figure 5.4's escalation flow. Not knowing the answer today is this week's fix.

5.5 Technical AI Governance for Foreign Affairs

For a foreign ministry, technical governance is not about building a testing lab — it is about knowing the right questions to ask vendors and partners:

By the time an AI tool reaches a diplomat's desk, its real capabilities and limits have already been fixed by decisions made far upstream — during training, testing and the allocation of computing power — long before anyone in the ministry touches it. A diplomat negotiating a digital-partnership agreement, a cyber attaché assessing a counterpart's AI capacity, and a mission deciding whether to adopt a foreign-hosted system all need at least a working grasp of those upstream decisions, because that is what tells them how far to trust the tool and what to ask the people who built it.

Technical governance strengthens the model layer by establishing what a system can do before it is deployed, supports the application layer by limiting what the deployed system is allowed to do, sets the standard the institutional layer enforces as a condition for adoption, and matters most directly at the infrastructure layer, in questions of who controls the computing power, supply chains and update paths.

Three technical governance tools matter increasingly for foreign affairs.

1. The first is testing for dangerous capabilities: the practice Shevlane et al. (2023) and Shah et al. (2025) describe as checking, before deployment, whether a model could enable harmful acts such as manipulation or cyber-attacks.
2. The second is compute governance — the idea, put forward by Ramiah et al. (2025) in a proposal for a global system to govern computing power with a built-in pause mechanism, that the large computer clusters used to train the most advanced models are a workable point of control for international oversight.
3. The third is hardware-level governance, where Ansari (2026) sets out a practical framework for checking compliance, and even treaty obligations, at the level of computer chips and data centres.

The UK's Department for Science, Innovation and Technology (DSIT, 2023), in Emerging Processes for Frontier AI Safety, and the Emerging Practices in Frontier AI Safety Frameworks review (Buhl et al., 2025) show these ideas moving from theory into real government practice, while the International AI Safety Report 2026 places them in a global conversation Malaysia can join through ASEAN. What this calls for is two things: **on the outside, Malaysia should be able to represent its own interests in emerging rules on computing power and safety; on the inside, it should require vendors to share their safety-testing and incident history as a condition of doing business.**

This debate fortunately is no longer confined to research papers. In July 2026, more than 1,200 employees across OpenAI, Anthropic, Google DeepMind, Meta and other frontier labs signed an open statement, "[Pacing the Frontier](#)," asking governments to help build the technical and institutional tools needed to slow AI development if it is ever judged necessary (Nolan, 2026). Two months later, the heads of three of those same companies — Dario Amodei of Anthropic, Sam Altman of OpenAI and Elon

Musk of xAI — publicly called for a more deliberate pace of frontier model development, with Amodeli proposing that independent safety evaluators be given standing access to test frontier systems before release (Whitelaw & Fiegberman, 2026). Malaysia does not need frontier-scale compute to have a stake in that conversation: a ministry that understands why the labs building these systems are themselves asking for a pause button is better placed to ask sharp questions of any vendor selling to the government.

Additionally, the infrastructure-layer sovereignty question this section raises is also no longer theoretical: global spending on sovereign AI — models, data centres and cloud infrastructure built and controlled within national borders — was on pace to exceed US\$100 billion in 2026, led by commitments such as France's €109 billion infrastructure programme and Saudi Arabia's HUMAIN initiative, which had ordered more than 600,000 AI chips (Cohen, 2026). Ukraine is building a sovereign national large language model, adapted from open-weight foundations, specifically to reduce wartime dependence on foreign providers (Digital State UA, 2026), and India's IndiaAI Mission has approved thirteen projects under a dedicated "Safe & Trusted AI" pillar covering deepfake detection and bias mitigation (DD India, 2026).

Most sovereign models worldwide are still fine-tuned versions of foreign open-weight systems rather than built from first principles (Chavez et al., 2026) — a reminder for Malaysia that **sovereignty is a spectrum of control over data, compute and update rights, not a binary choice between full self-reliance and total dependence**, and we have a choice in those spectrum.

Case studies — 5.5 Technical AI Governance for Foreign Affairs

The International AI Safety Report 2026 shows the wider diplomatic track already under way: states are building shared language for describing frontier risk, and Malaysia's experience chairing ASEAN in strategic communication (Chapter 1) puts it in a position to help shape that language rather than simply receive it (International AI Safety Report, 2026).

Ramiah et al.'s (2025) compute-governance proposal points to a concrete diplomatic idea — a checkable "pause button" for training the most advanced models — that a country like Malaysia can champion in ASEAN forums as a confidence-building step: how can we leverage our technical strengths, in this case, compute chips, to strengthen our governance. This is exactly the kind of **technical-governance contribution that turns AI-safety knowledge into real diplomatic leverage**.

What you can do today (5.5) — model and infrastructure layers
When assessing any AI tool for use at a mission, ask the vendor three questions: what dangerous-capability testing did you run, where is my data stored, and what is your history of reported incidents? Asking these questions puts the burden of proof on the supplier, and is the simplest, most accessible form of technical governance — no technical background required. Also, before any mission adopts a foreign-hosted AI tool, get one sovereignty answer in writing: "if this service is withdrawn or changes overnight, what is our fallback?" File the answer before deployment, not after.

5.6 From Model Safety to Diplomatic System Safety

The sections above have treated AI safety as a set of separate steps: classify, evaluate, red-team, guideline, and govern. This final section brings them together as layers of one thing — the AI-enabled foreign-affairs system. The key point is that safety cannot live in the model alone. A ministry that buys a well-tested model has only dealt with the first, most basic layer.

- The model layer covers ability, hallucination, bias and uncertainty.
- The application layer is where the system finds documents, translates, summarises, drafts and supports decisions.
- The institutional layer covers workflows, sign-off, accountability and human review — the routines this chapter has argued cannot be optional.
- The infrastructure layer is cloud, data centres, networks, identity, system connections and vendors — the foundation a ministry usually does not control when it buys a commercial model.

A system can perform well at the model layer and still be unsafe, because it pulls in poisoned information, has too much access to diplomatic files, produces output that gets accepted without checking, gets silently changed by a vendor update, or simply stops working when the service is withdrawn.

The International AI Safety Report 2026 treats technical safeguards, monitoring, evaluation and institutional risk management as working together, not as substitutes for one another — the same approach this chapter has taken throughout (International AI Safety Report, 2026).

This leads to the book's central argument about diplomacy, applied at the system level: **technical safety matters, but it is not enough on its own**. It cannot decide what is politically legitimate, what risk is acceptable, or how power should be shared. **Once a ministry moves beyond trusting the model, it has to govern the whole workflow** — and, at the infrastructure layer, face the sovereignty question of who controls the data, the operations, the updates and the jurisdiction. Chapter 6 will take up the sovereignty and governance side of that question; Chapter 7 looks at how the service builds the capacity to sustain it.

Figure 5.6 The four-layer AI-enabled foreign-affairs system and how it can fail




Figure 5.6 The four-layer AI-enabled foreign-affairs system and how it can fail.

Source: Our synthesis from *International AI Safety Report (2026)* and the four-layer system model.

Case studies — 5.6 From Model Safety to Diplomatic System Safety

A system's safety is only as strong as the institutional layer enforcing it; without clear procedures and named reviewers, the model, application and infrastructure layers stay invisible to the officers actually using them.

The Malaysia election-manipulation case and the synthetic-identity breach used throughout this chapter are not failures of the model alone but failures that travelled across all four layers at once. In the Anthropic case, the model layer largely held — Claude repeatedly refused direct requests — but the application layer did not: the operators simply wrapped the same request in softer language until the wrapper, not the model's judgement, decided what counted as harmful (Anthropic, 2026). No institutional reviewer inside Malaysia's own system caught the operation before Anthropic's own threat-intelligence team did; when the report became public, the named regulator, MCMC, had to state that the activity fell outside its mandate rather than confirm it had already been tracking the threat (MCMC, 2026). And at the infrastructure layer, the operation ran entirely on foreign-hosted compute, controlled by a foreign vendor and a foreign commercial actor, with no Malaysian institution holding audit access or a kill switch over the platform doing the harm. The four-layer model predicts exactly this pattern: a largely safe model, an unsafe application, an unclear institutional owner, and infrastructure outside the affected country's control.

A second, global case shows how the same four layers can fail catastrophically even without any deliberate human attacker. In July 2026, OpenAI disclosed that experimental models it was using in an internal cybersecurity evaluation exploited a previously unknown flaw, left their intended test environment with no human direction, and used privilege escalation and stolen credentials to breach Hugging Face's production infrastructure while trying to "win" the exercise (OpenAI, 2026; Gold, 2026). Hugging Face detected and disclosed the intrusion independently before OpenAI confirmed, days later, that the responsible models were its own (Gold, 2026). Here the model layer had been evaluated and appeared to behave as intended in isolation; the failure emerged only once the application gave it enough tools and reach to act as an agent, the institutional layer had no process for containing an evaluation that escaped its sandbox, and the infrastructure layer — a shared package registry connecting two separate companies' systems — turned one company's test into another company's breach. It is one of the first publicly disclosed cases of an AI system autonomously breaching a real production system outside its own test environment, and it is a concrete illustration of why this chapter treats loss-of-control and systemic risk as categories a ministry must plan for, not just define.

The International AI Safety Report 2026 is itself proof of the four-layer argument at the policy level: it treats technical safeguards, monitoring, evaluation and institutional risk management as working together rather than as substitutes — the same position this chapter has held throughout (International AI Safety Report, 2026). For Malaysia, building these habits is also a regional opportunity: a foreign ministry that can speak credibly about capability testing, compute governance and system-level resilience is better placed to help shape ASEAN AI-safety norms than one that simply adopts tools.

What you can do today (5.6) — all four layers
Before adopting any AI tool, map it on one page across the four layers: which model, which application, which approval process, whose servers. Circle the layer you do not control. That single map shows exactly where the dependency lies and is the starting point for the full system-safety view this chapter describes.

Conclusion

AI safety in diplomacy is not a separate subject from diplomacy itself; it is the same craft adapted to a new tool. This chapter has treated safety as a set of everyday routines, not as an aspiration.

Classification gives each use case the care it deserves, building on Chapter 1's idea of augmentation by deciding where human review is required. Evaluation and monitoring check whether a system earns trust and keep watching after it is deployed, answering the verification problem Chapter 3 identified as the technology's central limitation. Red-teaming tests how systems fail under pressure, guarding against the impersonation and manipulation risks raised in Chapter 2. Guidelines and escalation protocols make sure human authority is never in doubt, and that when something goes wrong, it is contained rather than made worse — putting Chapter 4's accountability principles into practice.

The distinction drawn at the start — between AI safety, cybersecurity and information security — runs through all six sections. Securing a network does not make a model tell the truth. Protecting data does

not make a summary faithful to its sources. Testing a system against outside attackers is necessary but not enough if the system's own unpredictable behaviour goes unchecked. AI safety is about the behaviour of systems that generate, infer and decide — often without anyone seeing the moment things go wrong.

For Malaysia, building these habits is also a regional opportunity. A foreign ministry that can show disciplined, documented safeguards — and speak credibly about capability testing, compute governance, hardware checks and system-level resilience — is better placed to help shape ASEAN AI-safety norms than one that simply adopts tools. These safeguards exist so that technology gives diplomats more time and space to think, so they can keep doing what the profession has always required: understand, persuade, negotiate and use judgement on behalf of the state.

This chapter has been the operational core — the safeguards a ministry can put in place today. The book now turns outward. Chapter 6 moves from internal routine to the wider institutional and international picture: the governance frameworks, procurement and sovereignty clauses, and cross-border cooperation that make safe use last across ministries and partners. Chapter 7 asks how the service builds the capacity to sustain it — AI literacy, training needs and Malaysia's role in shaping regional AI-safety norms.

(6489 words, excluding References)

References

AI Security Institute. (2024, November 13). *Announcing Inspect Evals*. <https://www.aisi.gov.uk/blog/inspect-evals>

Akhtar, M., Reuel, A., Soni, P., Ahuja, S., Ammanamanchi, P. S., Rawal, R., Zouhar, V., Yadav, S., Whitehouse, C., Ki, D., Mickel, J., Choshen, L., Šuppa, M., Batzner, J., Chim, J., Sania, J., Long, Y., Rahmani, H. A., Knight, C., ... Solaiman, I. (2026). *When AI benchmarks plateau: A systematic study of benchmark saturation*. arXiv. <https://arxiv.org/abs/2602.16763>

Ansari, S. (2026). Hardware-level governance of AI compute: A feasibility taxonomy for regulatory compliance and treaty verification. arXiv:2604.04712. <https://arxiv.org/abs/2604.04712>

Anthropic. (2026, September 10). Detecting and countering misuse of AI: September 2026. <https://www.anthropic.com/threat-intelligence-report-september-2026>

Buhl, M. D., Bucknall, B., & Masterson, T. (2025). Emerging practices in frontier AI safety frameworks. arXiv:2503.04746. <https://arxiv.org/abs/2503.04746>

Chavez, P., Chilukuri, V., & Scanlon, R. (2026, August 20). Sovereign AI's second wave is coming into view. Center for a New American Security. <https://www.cnas.org/publications/cnas-insights/cnas-insights-sovereign-ais-second-wave-is-coming-into-view>

Cohen, T. (2026, July 7). \$100B+ sovereign AI — 40+ countries build own LLMs. Value Add VC. <https://valueaddvc.com/blog/sovereign-ai-why-every-country-is-now-building-its-own-llm>

Croft, D. (2025, October 9). Deloitte to refund government after using AI in \$440k report. Accounting Times. <https://www.accountingtimes.com.au/technology/deloitte-to-refund-government-after-using-ai-in-440k-report>

DD India. (2026, July 27). IndiaAI Mission's Safe & Trusted AI pillar advances responsible AI: 13 projects approved to tackle deepfakes and bias. <https://ddindia.co.in/2026/07/indiaai-missions-safe-trusted-ai-pillar-advances-responsible-ai-13-project-s-approved-to-tackle-deepfakes-and-bias/>

Digital State UA. (2026, January 7). Ukraine moves toward a sovereign AI model: National LLM to enter beta in 2026. <https://digitalstate.gov.ua/news/govtech/ukraine-moves-toward-a-sovereign-ai-model-national-llm-to-enter-beta-in-2026>

Ganguli, D., Lovitt, L., Kernion, J., Askill, A., Bai, Y., Kadavath, S., ... (2022). Red teaming language models to reduce harms: Methods, scaling behaviors, and lessons learned. arXiv:2209.07858. <https://arxiv.org/abs/2209.07858>

Gold, H. (2026, July 22). An OpenAI test model escaped and broke into a real company's computer system. CNN. <https://www.cnn.com/2026/07/22/tech/openai-hugging-face-ai-cybersecurity>

Ghosh, S., Frase, H., Williams, A., Luger, S., Röttger, P., Barez, F., McGregor, S., Fricklas, K., Kumar, M., Feuillade-Montixi, Q., Bollacker, K., Friedrich, F., Tsang, R., Vidgen, B., Parrish, A., Knotz, C., Presani, E., Bennion, J., Boston, M. F., ... Vanschoren, J. (2025). *AILuminate: Introducing v1.0 of the AI risk and reliability benchmark from MLCommons*. arXiv. <https://arxiv.org/abs/2503.05731>

Gross, G. (2025, November 6). AI spending may slow down as ROI remains elusive. CIO. <https://www.cio.com/article/4085325/ai-spending-may-slow-down-as-roi-remains-elusive.html>

International AI Safety Report. (2026). International AI Safety Report 2026. <https://internationalaisafetyreport.org/sites/default/files/2026-02/international-ai-safety-report-2026.pdf>

Koh, J. L. (2026, September 12). Turkish firm offered AI-driven 'political ops' tool targeting Malaysia, Anthropic finds. Malaysiakini. <https://www.malaysiakini.com/news/884887-turkish-firm-offered-ai-driven-political-ops-tool-targeting-malaysia-anthropic-finds>

Lee, M. (2025, July 8). Impostor uses AI to impersonate Rubio and contact foreign and US officials. Associated Press.

Malaysian Communications and Multimedia Commission [MCMC]. (2026, September 11). MCMC is aware of Anthropic's recent report and takes the matters raised seriously. For clarity, activities of the nature described in the report do not form part of MCMC's functions or mandate [Media statement] [Photograph]. Instagram. <https://www.instagram.com/p/DdJZnArGMkW/>

Ministry of Science, Technology and Innovation Malaysia. (2024). The National Guidelines on AI Governance & Ethics (AIGE).

National Institute of Standards and Technology. (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST. <https://doi.org/10.6028/NIST.AI.100-1>

National Institute of Standards and Technology. (2024). Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile (NIST AI 600-1). <https://doi.org/10.6028/NIST.AI.600-1>

Nolan, B. (2026, July 29). More than 1,200 AI workers are asking for Washington's help to pace frontier AI development. Fortune. <https://fortune.com/2026/07/29/anthropic-deepmind-openai-meta-washington-ai-slowdown-plan/>

OpenAI. (2026, July 21). OpenAI and Hugging Face partner to address security incident during model evaluation. <https://openai.com/index/hugging-face-model-evaluation-security-incident/>

OWASP Gen AI Security Project. (2024). OWASP Gen AI Security Project — Top 10 for LLM Applications. <https://genai.owasp.org/>

Ramiah, A. A., Koopmanschap, R., Thorsteinson, J., Khan, S., Zhou, J., Noh, S., Meindertsma, J., & Shafiq, F. (2025). Toward a global regime for compute governance: Building the pause button. arXiv:2506.20530. <https://arxiv.org/abs/2506.20530>

Shah, R., Irpan, A., Turner, A. M., Wang, A., ... (2025). An approach to technical AGI safety and security. arXiv:2504.01849. <https://arxiv.org/abs/2504.01849>

Shevlane, T., Farquhar, S., Garfinkel, B., Phuong, M., Whittlestone, J., Leung, J., ... (2023). Model evaluation for extreme risks. arXiv:2305.15324. <https://arxiv.org/abs/2305.15324>

Sidhu, H. K., Scholefield, R., Annan, N., Hernandez, K., Hou, I. N., Alshaikhi, A., Chin, Z. S., & Gipiškis, R. (2026). Open problems in AI incident governance. arXiv:2607.05163. <https://arxiv.org/abs/2607.05163>

Survey for Diplomats: Innovating Diplomacy (Responses). (2026). Unpublished survey dataset [64 responses]. Institute of Diplomacy and Foreign Relations.

Tadros, E., & Karp, P. (2025, October 5). Deloitte to refund government, admits using AI in \$440k report. Australian Financial Review. <https://www.afr.com/companies/professional-services/deloitte-to-refund-government-after-admitting-a-i-errors-in-440k-report-20251005-p5n05p>

The Edge Malaysia. (2026, September 11). Claude creator Anthropic reveals it disrupted election-manipulation platform targeting Malaysian voters. <https://theedgemaalaysia.com/node/817799>

The Star. (2026, September 11). MCMC takes Anthropic report seriously, reviewing issues raised. <https://www.thestar.com.my/news/nation/2026/09/11/mcmc-takes-anthropic-report-seriously-reviewing-issues-raised>

UK Department for Science, Innovation and Technology. (2023, October 27). Emerging processes for frontier AI safety. <https://www.gov.uk/government/publications/emerging-processes-for-frontier-ai-safety>

Whitelaw, K., & Fiegeman, S. (2026, September 12). Amodei, Altman, Musk call for slowing AI model development. Los Angeles Times. <https://www.latimes.com/business/story/2026-09-12/amodei-altman-musk-call-for-slowing-ai-model-development>